

docs4llm

Documentation that LLMs can use.

Turn any docs URL into a hosted MCP server - so agents can search, understand, and act.



01

DOCS URL
AI ANALYSIS
MCP ENDPOINT

Team FireCrackers
team-6F2C797C119E

Maharaja Surajmal Institute of Technology

The problem is not missing documentation.

It is documentation that an agent cannot reliably use.

Teams are forced to translate human-readable pages into agent-readable context again and again.

Affected

Developers / Product teams / AI agents

Current choices: static docs, RAG chatbots, or hand-built MCP wrappers.

What agents receive today

- **Docs site**
pages
- **OpenAPI spec**
schema
- **GitHub README**
repository
- **Changelog**
history


What they need instead

A dependable tool interface.

One URL.

A usable agent interface.


docs4llm transforms documentation into a hosted MCP endpoint -
not another layer of text for an LLM to guess through.



Why it is better: structured tools let agents perform the documented work instead of merely summarizing it.

The architecture is simple because the job is specific.

Mermaid.js system flow



Convex Integration in progress: project state, workflow events, and realtime conversion status.

Current stack / Next.js + React / TypeScript / Gemini / Supabase + Drizzle / Upstash QStash / MCP SDK

Mermaid: DocsURL --> Ingestion --> Gemini --> MCPTools --> AIClientes

Four parts. One useful outcome.

Make the documentation callable.

01 Ingest

Docs sites, OpenAPI, Postman, GitHub, MD/MDX, and HTML.

02 Understand

Gemini extracts endpoints, workflows, and documentation meaning.

03 Generate

Create typed MCP tools and a documentation agent.

04 Validate

Check tool schemas and smoke-test the runtime before publishing.

The innovation is not another chat window. It is a reusable tool layer for every AI client.

Who benefits?

Developers

Less hunting. Less manual setup.

Product teams

One source of truth across tools.

AI agents

Structured context they can call.

Feasible from day one.

01 Serverless deployment

Next.js and hosted MCP endpoints keep the operating model lightweight.

02 Queue-based processing

Crawling and AI analysis scale independently from live tool calls.

03 Convex-ready backend

Realtime state and workflow coordination are being integrated for the Convex track.

Build once. Let every agent read the docs.

docs4llm turns the web's documentation into dependable, reusable AI infrastructure.

Prototype demo

- Paste a docs URL
- Watch crawl and analysis
- Receive a hosted MCP endpoint
- Connect it to an AI client

Live demo: building phase

mcp.json

```
{  
  "mcpServers": {  
    "docs4llm": {  
      "url": ".../mcp"  
    }  
  }  
}
```

[Why this should be selected](#)

A practical tool layer for
AI-assisted builders.